

TENSORFLOW AND DEEP LEARNING

Lili Jiang, Xuan-Son Vu , Michele Persiani

Department of Computing Science

Organized by HPC2N

8-9 May, 2019



UMEÅ UNIVERSITY

DAY 1-MORNING: TENSORFLOW

- 9:00-10:00 Lecture: Introduction to Tensorflow.
- 10:00-10:30 Exercise: Play with basic Tensorflow and TensorBoard.

10:30 Coffee break

- 11:00-12:00 Exercise: Play with Tensorflow on real dataset.

12:00-13:00 Lunch



UMEÅ UNIVERSITY

DAY 1-AFTERNOON: DEEP LEARNING MODELS IN NLP-1

- 13:00-14:00 Lecture: Introduction of Deep Learning on Natural Language Processing (NLP).
- 14:00-14:30 Exercise: Run data representation in basic.

14:30 Coffee break

- 15:00-16:00 Exercise: Play with NLP-based deep learning models (RNN, Transformer).



UMEÅ UNIVERSITY

DAY 2-MORNING: DEEP LEARNING MODELS IN NLP-2

- 9:00-10:00 Lecture: Introduction of different state-of-the-art deep learning models on NLP.

10:15 Coffee break

- 10:30-11:45 Exercise: Play with NLP-based deep learning models (Bert and ELMo).

- 11:45-12:00 Lecture: Wrap-up

12:00-13:00 Lunch



UMEÅ UNIVERSITY

DAY 2-AFTERNOON: GAN (GENERATIVE ADVERSARIAL NETWORKS)

- 13:00-14:00 Lecture: Introduction of GAN
 - 14:00-15:15 Exercise: Play with GAN.
- 15:15 Coffee break**
- 15:45-16:15 Lecture: Wrap-up



UMEÅ UNIVERSITY

COURSE WEBSITES

- Course github

<https://github.com/ddmatumu/TFnDeepLearning>

- HPC2N website

<https://www.hpc2n.umu.se/events/courses/deep-learning-spring-2019>

- Feedback:

www.menti.com (Code: 21 40 80)



UMEÅ UNIVERSITY

- Finished the installation by following the instructions on github?
- Access internet (for those from other Uni)?
- How many Windows users?



UMEÅ UNIVERSITY

INTRODUCTION TO TENSORFLOW

Lili Jiang, Xuan-Son Vu , Michele Persiani

Organised by HPC2N

08 May, 2019



UMEÅ UNIVERSITY

TENSORFLOW

- An open source library for numerical computation.
- Launched by Google in 2015, November
- Tensorflow 2.0 Alpha released in 2019, March



UMEÅ UNIVERSITY

WHY TENSORFLOW

- Flexibility + Scalability



WHY TENSORFLOW

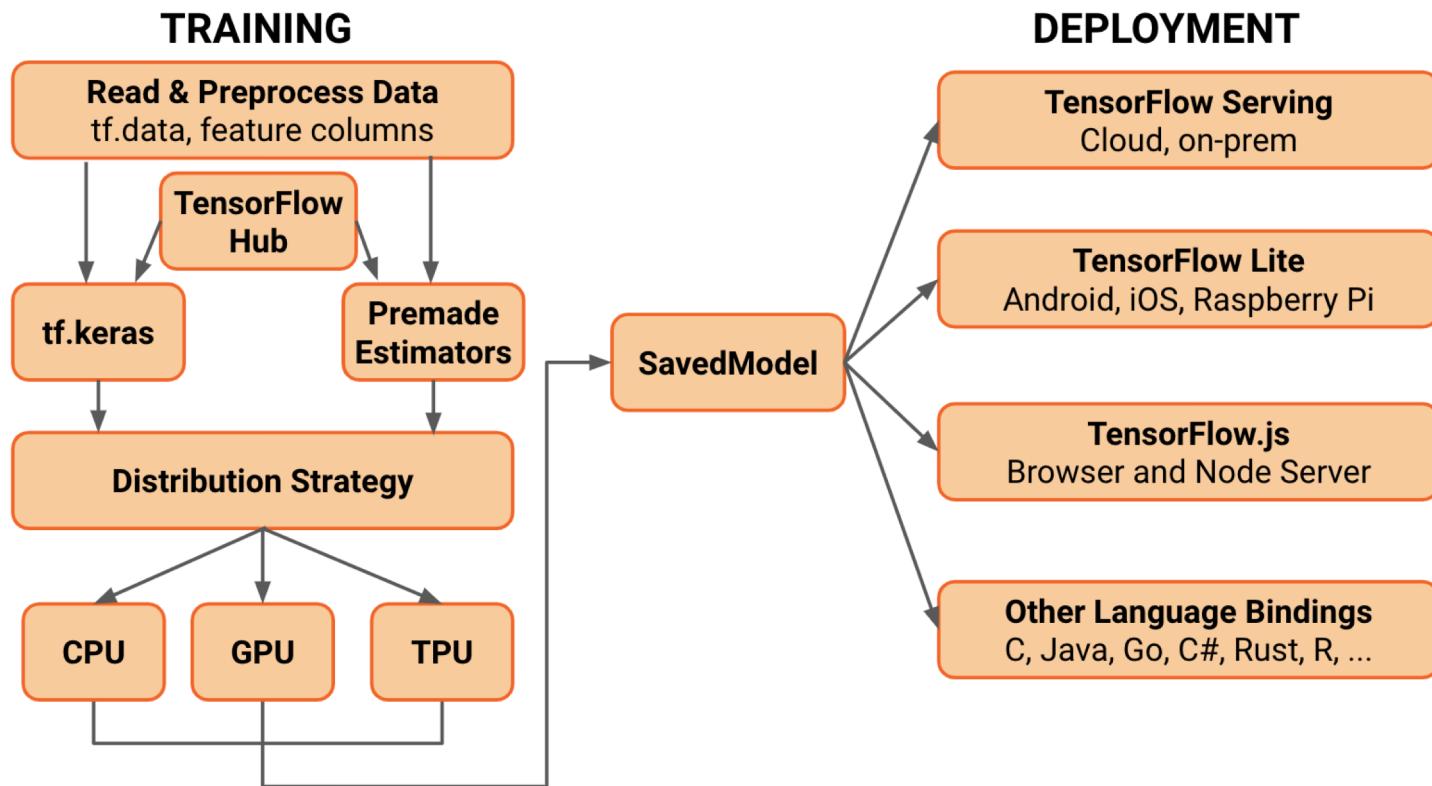
- Chainer and PyTorch
 - Flexible enough for research but less scalable
- Caffe and MXNet
 - Scalable but less flexible
- Keras
 - Simpler/beginner, but difficult for customization.
- Theano
 - High performance, but difficult syntax



UMEÅ UNIVERSITY

WHY TENSORFLOW

- And last but not least:
 - Tensorflow meets industry's requirements



WHAT IS TENSOR

- **An n-dimensional array**
- 0-d tensor: scalar (number)
- 1-d tensor: vector
- 2-d tensor: matrix
-



UMEÅ UNIVERSITY

WHAT IS TENSOR

- $t_0 = 3$ #scalar is 0-d tensor
- $t_1 = ["cat", "dog",]$ # 1-d tensors
- $t_1 = [[1, 2], [2, 3]]$ # 2-d tensors
-



UMEÅ UNIVERSITY

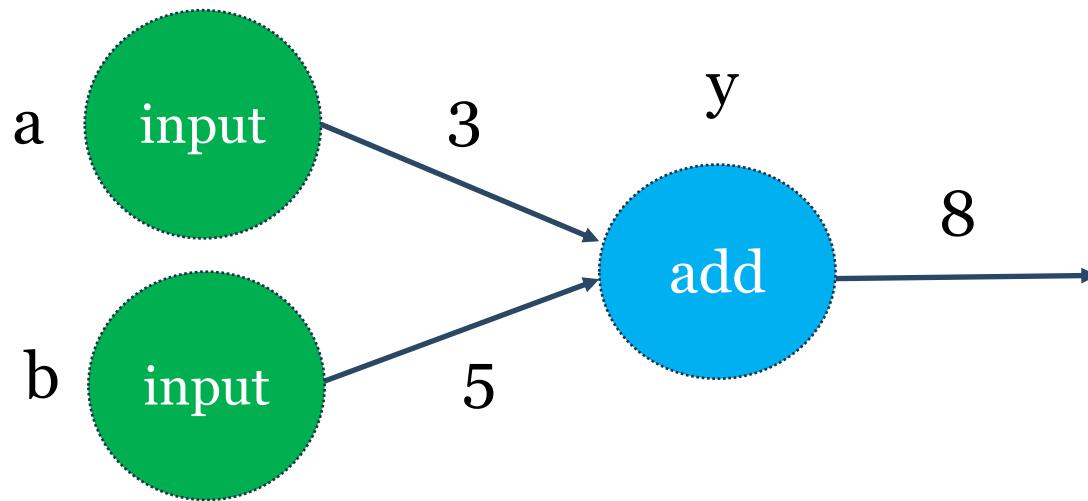
Tensorflow = Tensor + Flow = Data + Flow



UMEÅ UNIVERSITY

DATA FLOW GRAPH

- Use a session to execute operations in the graph



Nodes: operators, variables, and constants

Edges: tensors

TENSORFLOW

- Each node is an operation (op).
- Data is represented as Tensors.
Op takes Tensors and returns Tensors.
- Variables maintain state across executions of the graph.



UMEÅ UNIVERSITY

TENSORFLOW "HELLO WORLD"

```
import tensorflow as tf  
a = tf.constant(3)  
b = tf.constant(5)  
y = tf.add(a, b)  
  
with tf.Session() as sess:  
    print(sess.run(x))
```

Graph and sessions

TWO PHASES IN THE PROGRAM

- Construct the computation graph.
- Executes a graph in the context of a Session.



UMEÅ UNIVERSITY

TF.SESSION()

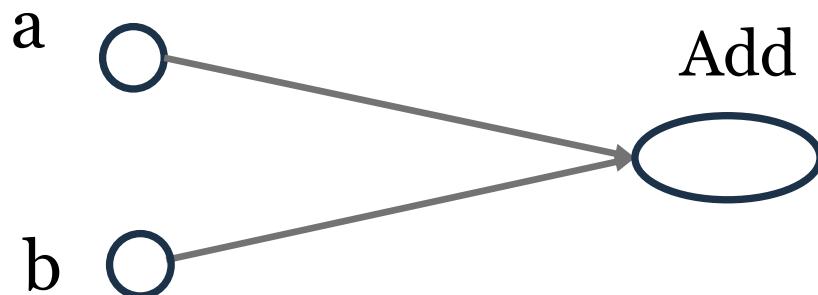
- A Session object encapsulates the environment in which Operation objects are executed, and Tensor objects are evaluated.
- Session will also allocate memory to store the current values of variables.



UMEÅ UNIVERSITY

VISUALIZE IN TENSORBOARD

- Take **Add** as an example:



UMEÅ UNIVERSITY

TENSORBOARD

- Visualize the computation graph
- Present the vector space, to see how they related with each other.
- Help tuning / training parameters
- Debug your computational graphs



UMEÅ UNIVERSITY

TENSORFLOW DATA TYPES

- Constant
 - `a = tf.constant(2, name='a')`
- Variable
 - `tf.Variable(3, name="scalar")`
- Placeholder
 - E.g., function $f(x, y) = 2 * x + y$ without knowing value of x or y.
 - x, y are placeholders for the actual values.



UMEÅ UNIVERSITY

CONSTANT

```
import tensorflow as tf  
a = tf.constant(2, name = 'a')  
b = tf.constant([2, 3], name = 'b')  
y = tf.add(a, b, name = 'add')  
with tf.Session() as sess:  
    print(sess.run(y))
```



UMEÅ UNIVERSITY

RANDOM CONSTANT

- tf.random_normal
- tf.truncated_normal
- tf.random_uniform
- tf.random_shuffle
- tf.random_crop
- tf.multinomial
- tf.random_gamma



UMEÅ UNIVERSITY

TIP OF USING CONSTANT

- Only use constants for primitive types.
- Use variables or readers for more data that requires more memory



UMEÅ UNIVERSITY

VARIABLE

- # create variables with `tf.Variable`
- `s = tf.Variable(2, name="scalar")`
- `m = tf.Variable([[0, 1], [2, 3]], name="matrix")`
- `W = tf.Variable(tf.zeros([784, 10]))`



UMEÅ UNIVERSITY

VARIABLE

```
# create variables with tf.get_variable
```

- `s = tf.get_variable("scalar",
initializer=tf.constant(2))`
- `m = tf.get_variable("matrix",
initializer=tf.constant([[1, 2], [2, 3]]))`
- `W = tf.get_variable("matrix", shape=(10, 10),
initializer=tf.zeros_initializer())`

(**recommended!**)



UMEÅ UNIVERSITY

VARIABLE

- Use variables for more data requiring memory.
- tf.variable or tf. Variable?
- Tf. constant is an op
- tf. Variable is a class.



UMEÅ UNIVERSITY

ARITHMETIC OPS

- tf.abs
- tf.negative
- tf.square
- tf.round
- tf.sqrt
- tf.exp



UMEÅ UNIVERSITY

TIP OF USING VARIABLE

#initialization

```
W = tf.Variable(tf.zeros([784,10]))
```

with tf.Session() as sess:

```
    sess.run(W.initializer)
```



UMEÅ UNIVERSITY

TIP OF USING VARIABLE

#initialize all variables at once:

with tf.Session() as sess:

```
sess.run(tf.global_variables_initializer())
```

Initialize only a subset of variables:

with tf.Session() as sess:

```
sess.run(tf.variables_initializer([a, b]))
```



UMEÅ UNIVERSITY

PLACEHOLDER

- For example, a function $f(x, y) = 2 * x + y$ without knowing value of x or y.
- x, y are placeholders for the actual values.



UMEÅ UNIVERSITY

PLACEHOLDER (EXAMPLES)

```
tf.placeholder(dtype, shape=None, name=None)
```

- a = tf.placeholder(tf.float32, shape=[3])
- b = tf.constant([5, 5, 5], tf.float32)



UMEÅ UNIVERSITY

PLACEHOLDER VS. VARIABLE

- Variable defines: shape + its values + data type
- Placeholder defines: shape + data type

- Variable: hold variables to be trained
- Placeholder: hold data to feed into the model

- `b = tf.Variable(tf.zeros([1]))`
- `x = tf.placeholder(tf.float64, [None, 25])`



UMEÅ UNIVERSITY

TENSORFLOW AND NUMPY

tensor = tf.constant(np_array)

Example:

```
sess = tf.Session()
```

```
a = tf.zeros([2, 3], np.int32)
```

```
a_out = sess.run(a)
```

```
print(type(a_output)) # => <class 'numpy.ndarray'>
```



UMEÅ UNIVERSITY

CONTROL DEPENDENCIES

```
tf.Graph.control_dependencies(control_inputs)
```

defines which ops should be run first

your graph g have 5 ops: a, b, c, d, e

```
g = tf.get_default_graph()
```

```
with g.control_dependencies([a, b, c]):
```



UMEÅ UNIVERSITY

THE TRAP OF LAZY LOADING

```
x = tf.Variable(10, name='x')
```

```
y = tf.Variable(20, name='y')
```

with tf.Session() as sess:

```
    sess.run(tf.global_variables_initializer())
```

```
    for _ in range(10):
```

```
        sess.run(tf.add(x, y))
```

```
writer.close()
```



UMEÅ UNIVERSITY

NORMAL LOADING

```
x = tf.Variable(10, name='x')
y = tf.Variable(20, name='y')
z = tf.add(x, y)
```

with tf.Session() as sess:

```
    sess.run(tf.global_variables_initializer())
```

```
    for _ in range(10):
```

```
        sess.run(z)
```

```
    writer.close()
```



UMEÅ UNIVERSITY

TIP / WARNING

- Separate definition of ops from computing/running ops
- Ensure that function is also loaded once the first time it is called.



UMEÅ UNIVERSITY

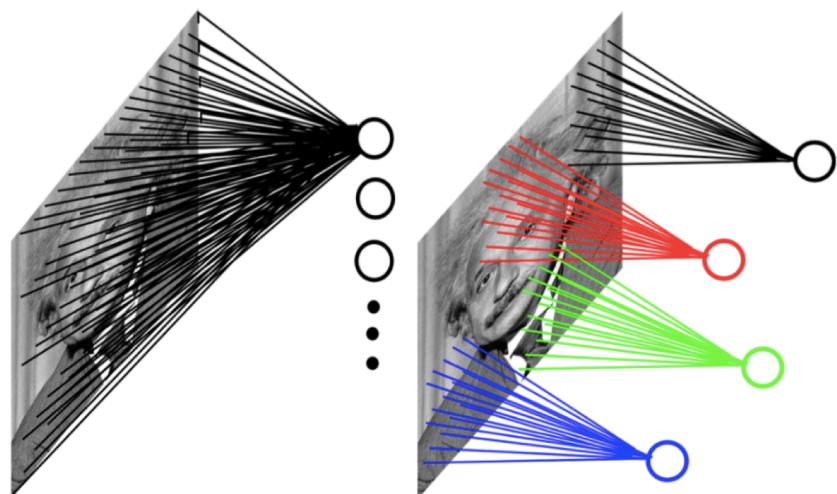
Instead of $a+b$,
build neural networks in tensorflow?
CNN? RNN?



UMEÅ UNIVERSITY

CONVOLUTIONAL NEURAL NETWORK

- CNNs are a special type of neural network **whose hidden units are only connected to local receptive field**. The number of parameters needed by CNNs is much smaller.

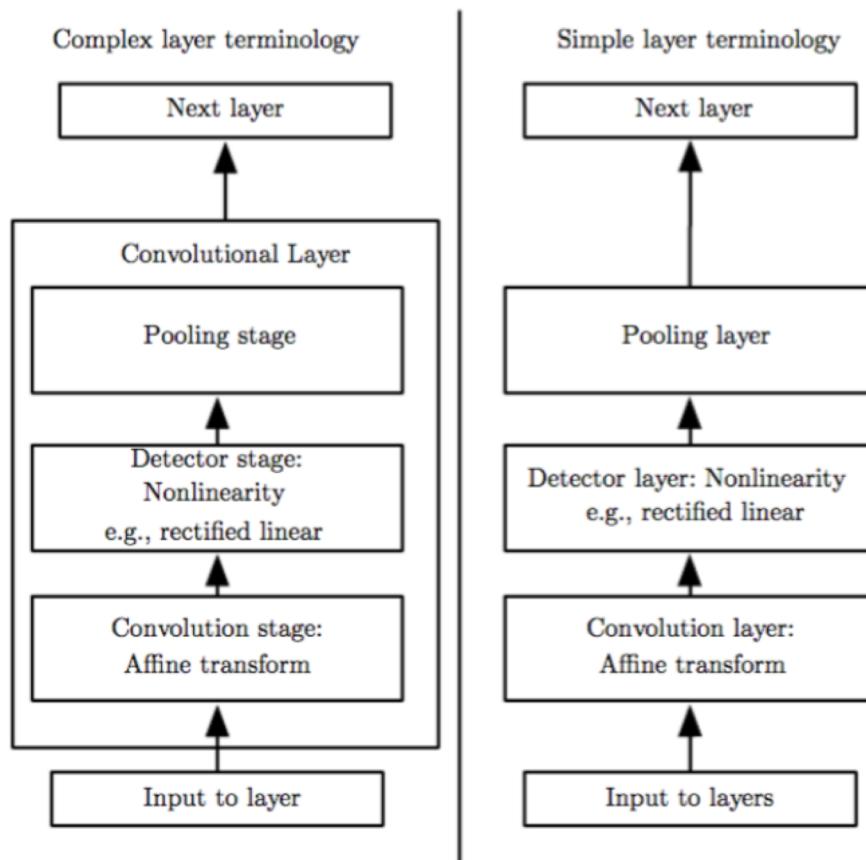


Example: 200x200 image

a) fully connected: 40,000
hidden units => 1.6 billion
parameters

b) CNN: 5x5 kernel, 100 feature
maps => 2,500 parameters

THREE STAGES OF A CONVOLUTIONAL LAYER

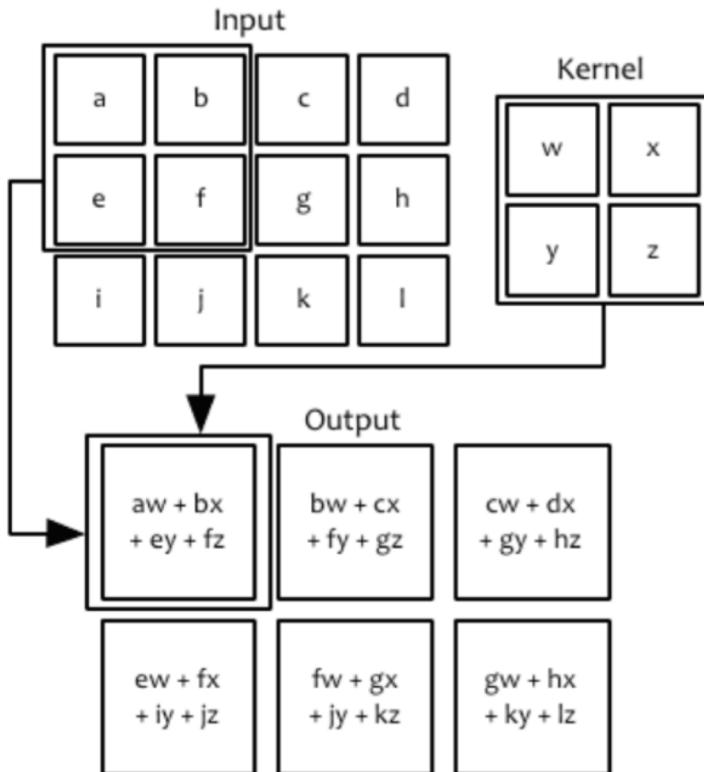


Convolution stage

Nonlinearity: a nonlinear transform such as rectified linear or tanh

Pooling: output a summary statistics of local input, such as max pooling and average pooling

CONVOLUTION OPERATION IN CNN

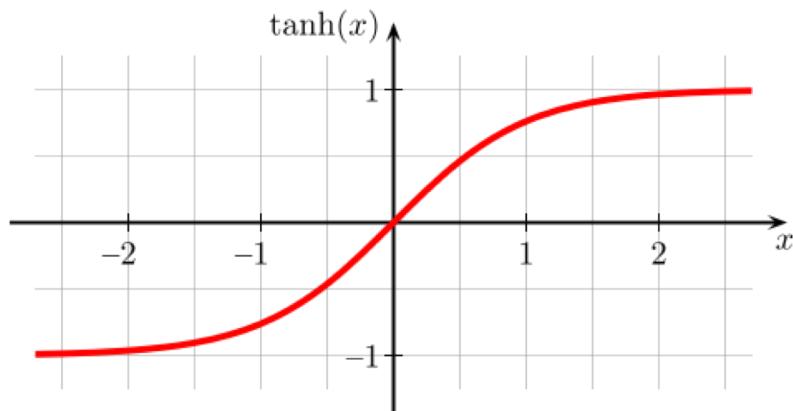


- Input: an image (2-D array) x
- Convolution kernel/operator(2-D array of learnable parameters): \mathbf{w}
- Feature map (2-D array of processed data): \mathbf{s}
- Convolution operation in 2-D domains

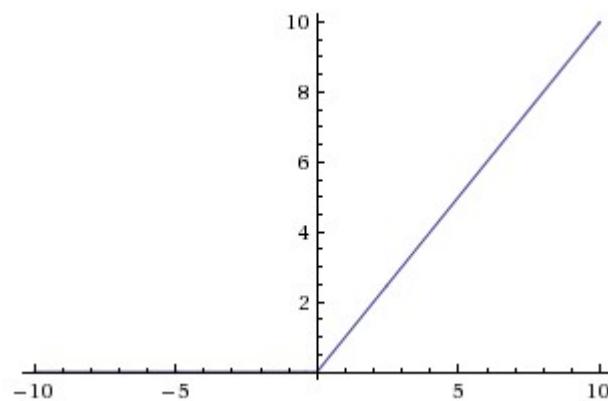
$$s[i, j] = (x * w)[i, j] = \sum_{m=-M}^M \sum_{n=-N}^N x[i + m, j + n] w[m, n]$$

NON-LINEARITY

Tanh(x)



ReLU

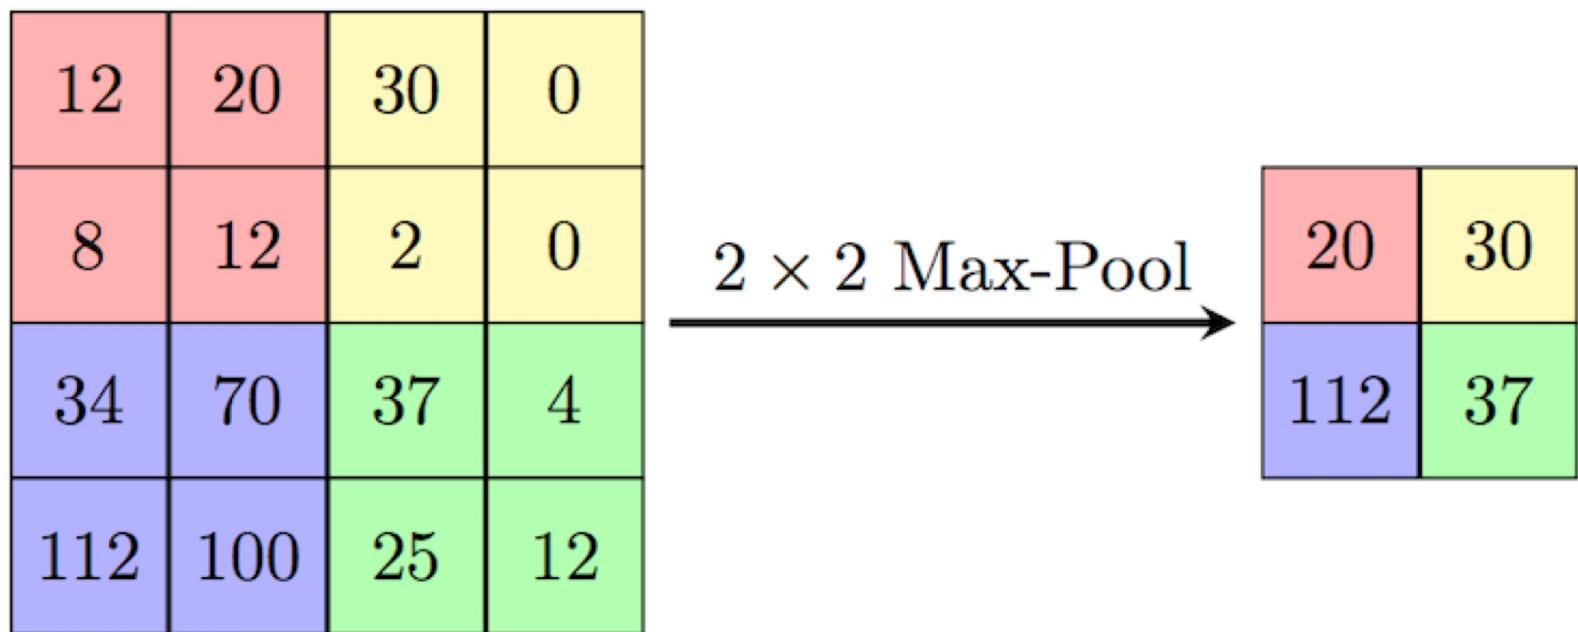


$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$f(x) = \max(0, x)$$

POOLING

- Common pooling operations:
 - **Max pooling**: reports the maximum output within a rectangular neighborhood.
 - **Average pooling**: reports the average output of a rectangular neighborhood (possibly weighted by the distance from the central pixel).



CNN IN TENSORFLOW

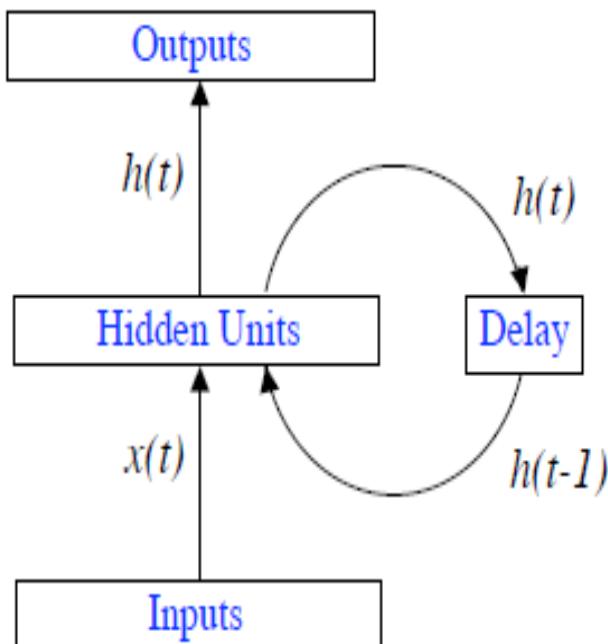
- tf.layers module -> construct a neural network
 - conv2d(): arguments - number of filters, filter kernel size, padding, and activation function.
 - max_pooling2d(): arguments - pooling filter size and stride.
 - dense(): arguments - number of neurons and activation function.



UMEÅ UNIVERSITY

WHAT ARE RNNs?

The simplest form of *fully recurrent neural network* with the previous set of hidden unit activations feeding back into the network along with the inputs



$$h(t) = f_H(W_{IH}x(t) + W_{HH}h(t - 1))$$

$$y(t) = f_O(W_{HO}h(t))$$

f_H and f_O are the activation function for hidden and output unit; W_{IH} , W_{HH} , and W_{HO} are connection weight matrices which are learnt by training.

TRAINING RNNS

Back Propagation Through Time (BPTT), an extension of the back-propagation (BP), is often used to learn RNNs.

- The output of this RNN is \hat{y}_t

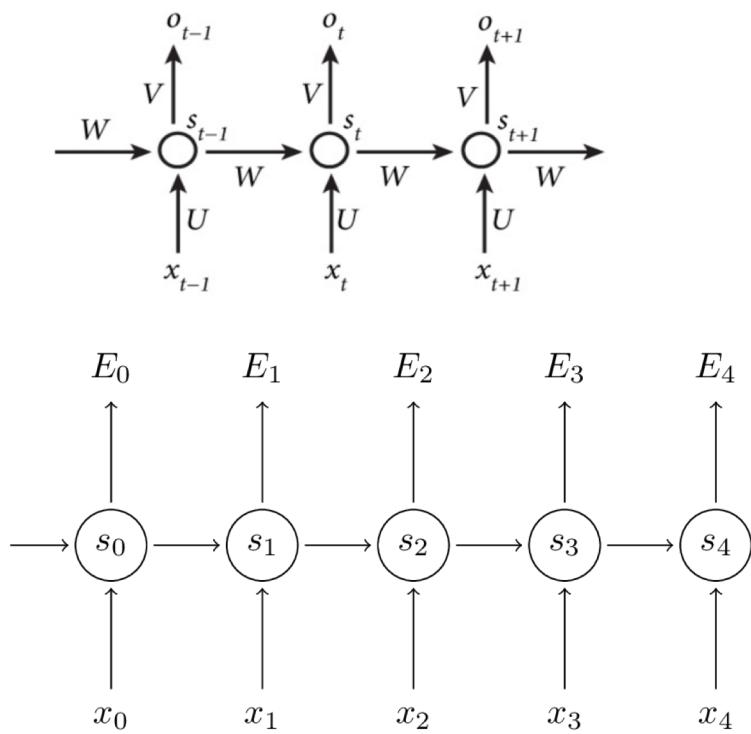
$$s_t = \tanh(Ux_t + Ws_{t-1})$$

$$\hat{y}_t = \text{softmax}(Vs_t)$$

- The loss/error function:

$$E_t(y_t, \hat{y}_t) = -y_t \log \hat{y}_t$$

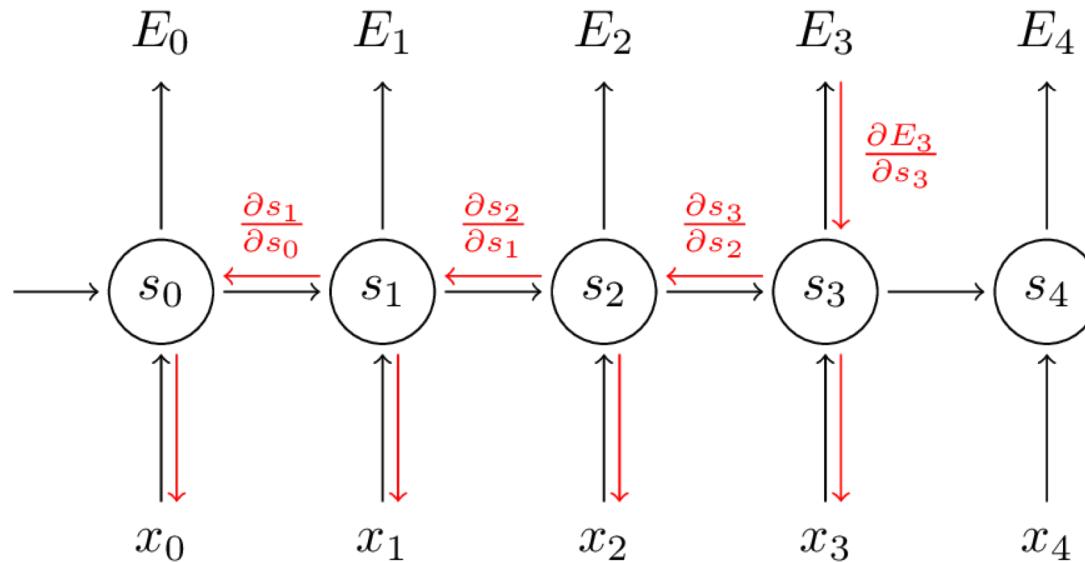
$$E(y, \hat{y}) = \sum_t E_t(y_t, \hat{y}_t)$$



TRAINING RNNS

$$\frac{\partial E_3}{\partial W} = \sum_{k=0}^3 \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_k} \frac{\partial s_k}{\partial W}$$

Because W is used in every step up to the output we care about, we need to back-propagate gradients from $t = 3$ through the network all the way to $t = 0$



RNN IN TENSORFLOW

`tf.Graph() + tf.Session()`

- Initialize placeholders which are filled with batches of training data.
- Define the RNN model and to calculate the output values (logits) -> loss value
- Use an Optimizer to optimize the weights of the RNN.



UMEÅ UNIVERSITY

RNN FLAVOR IN TENSORFLOW

Cells

- `tf.nn.rnn_cell.BasicRNNCell(num_hidden)`
- `tf.nn.rnn_cell.BasicLSTMCell(num_hidden)`
- `GruCell, MultiRNNCell`

Container

- `tf.nn.static_rnn(cell, data, dtype=tf.float32)`
- `dynamic_rnn (cell, data,)`
- `static_bidirectional_rnn(.....)`

RNN FLAVOR IN TENSORFLOW

Cells

`tf.keras.layers.SimpleRNNCell`
`tf.keras.layers.LSTMCell`

- `tf.nn.rnn_cell.BasicRNNCell(num_hidden)`
- `tf.nn.rnn_cell.BasicLSTMCell(num_hidden)`
- `GruCell`, `MultiRNNCell`

Container

- `tf.nn.static_rnn(cell, data, dtype=tf.float32)`
- `dynamic_rnn (cell, data,)`
- `static_bidirectional_rnn(....)`
 `keras.layers.RNN(cell)`
 `keras.layers.RNN(cell, unroll=True)`
 `keras.layers.Bidirectional(keras.layers.RNN(cell, unroll=True))`

EXCERCISE 1

- Basic operations in Tensorflow
- Linear Regression using tensorflow



UMEÅ UNIVERSITY

EXCERCISE 2

- Predicting HPC total energy consumption using tensorflow



UMEÅ UNIVERSITY

FEEDBACK ALWAYS WELCOME !

- Go to www.menti.com
- Code: 21 40 80



UMEÅ UNIVERSITY

Thanks



UMEÅ UNIVERSITY